

Pr. Didier Donsez ([didier.donsez@imag.fr](mailto:didier.donsez@imag.fr))

Pr. Olivier Gruber ([olivier.gruber@imag.fr](mailto:olivier.gruber@imag.fr))

Laboratoire d'Informatique de Grenoble  
Université de Grenoble-Alpes

# This Course – Two Parts

- Part One – Pr. Olivier Gruber – 50%
  - 5 weeks, widening your horizon as developers
  - Wandering outside of the Linux world
  - Also a primer before part two
- Part Two – Pr. Didier Donsez – 50%
  - 5 weeks, experimenting with an IoT infrastructure
  - Sensors → Gateway → Data-gathering Server



# Machine to Machine

Freely-adapted excerpts from Wikipedia:

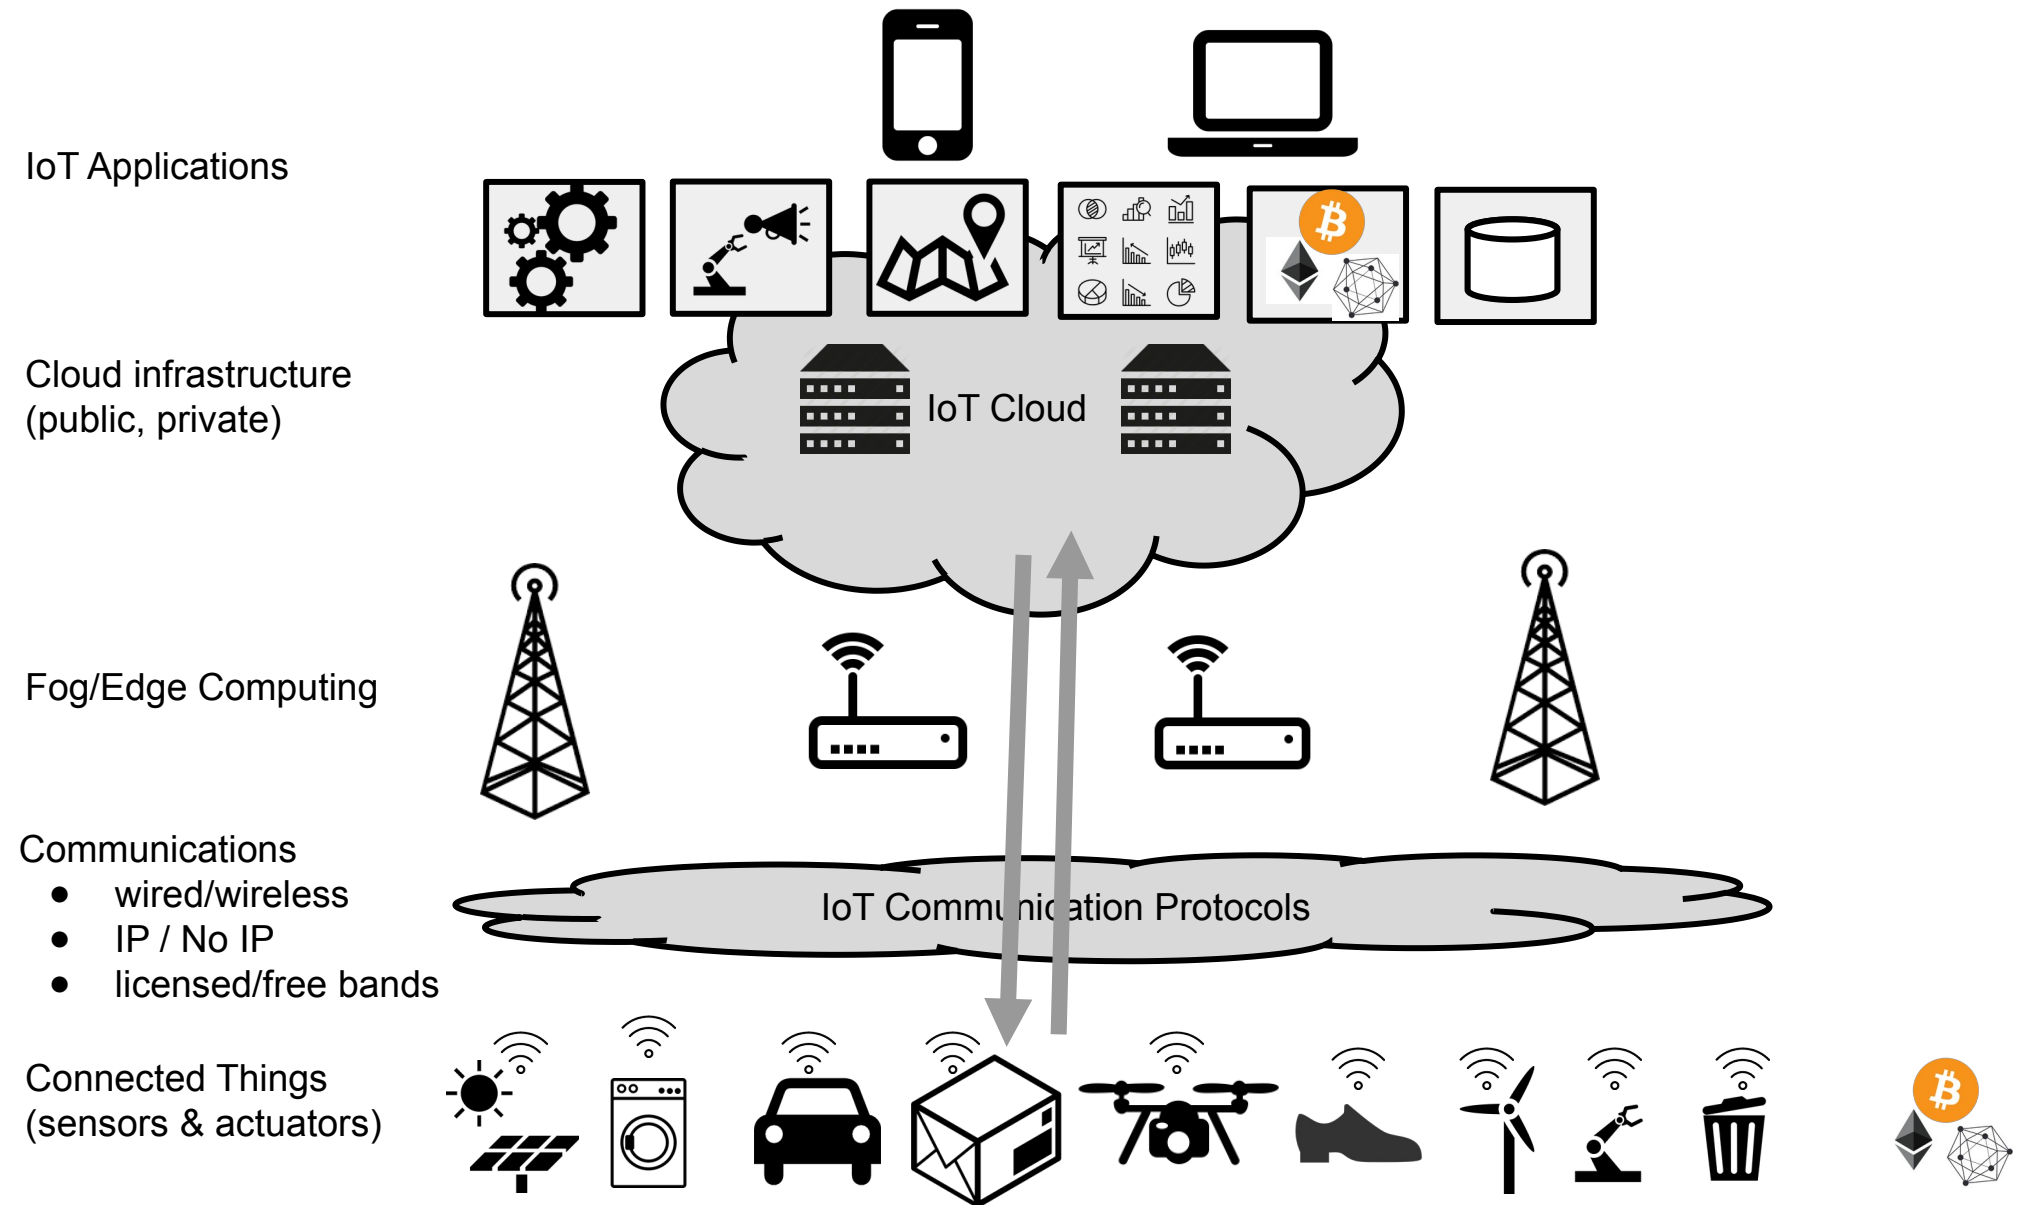
*Machine to machine (M2M) is direct communication between devices using any communications channel, including wired and wireless.*

*M2M communication can include industrial instrumentation, enabling sensors to communicate the information it records to software systems that can use it.*

*The Internet of things (IoT) describes the network of physical objects—“things”—that embed sensors, software, and other technologies for the purpose of connecting and exchanging data with other devices and software systems over the Internet.*

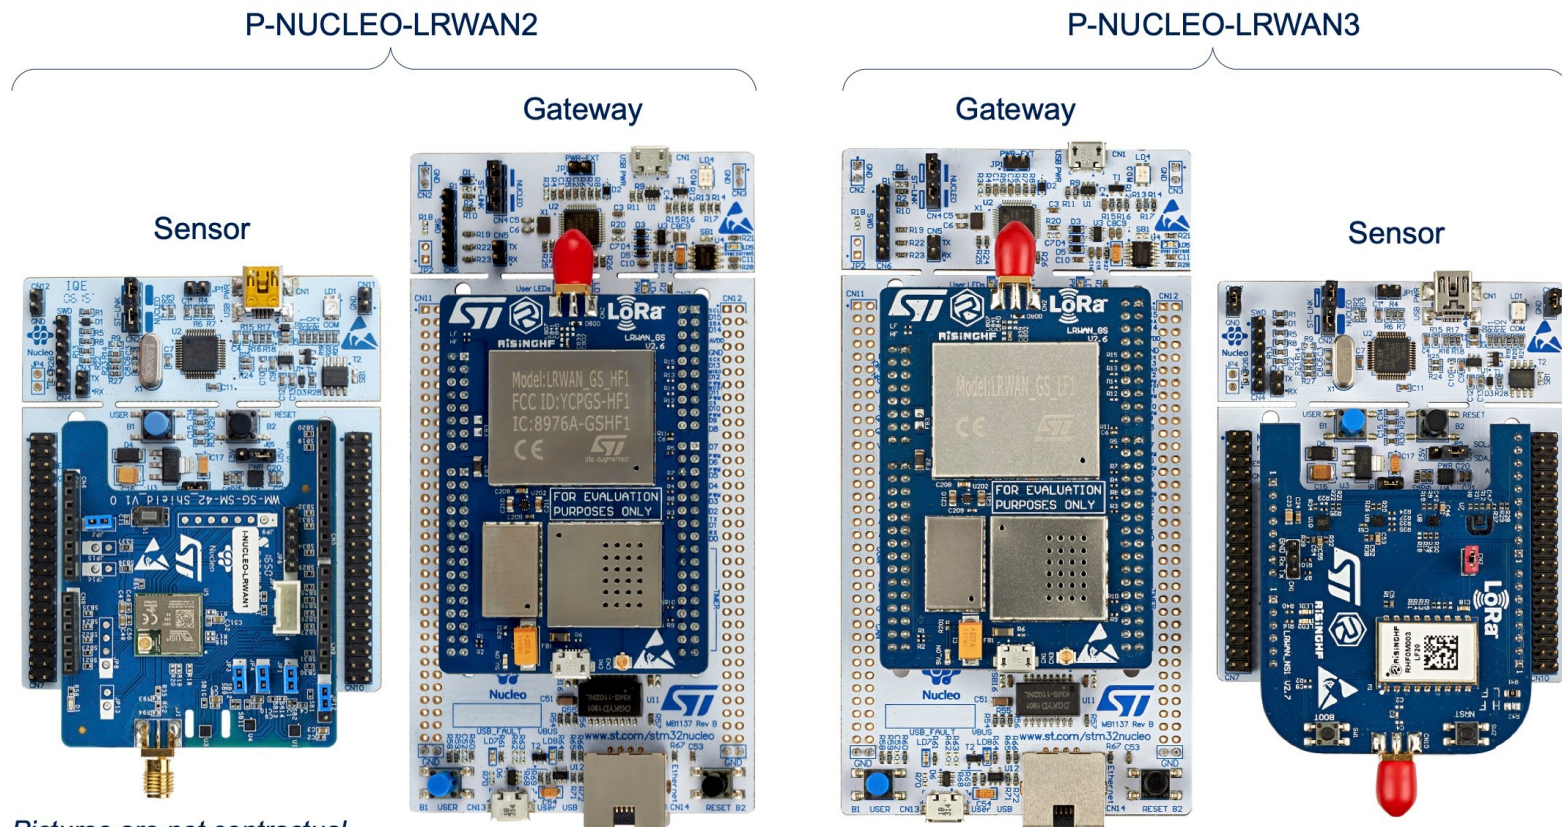
*Traditional fields of embedded systems, wired or wireless sensor networks, control systems, automation (including home and building automation), and many others all contribute to enabling the Internet of things.*

# IoT reference architecture



# P-Nucleo-LRWAN Starter Packs

**ToDo:** Pick up your P-Nucleo-LRWAN2 at the fablab, after January 30th, but before the lectures start with Pr. Didier Donsez.



# Software – RIOT

The friendly operating system for the internet of things.

RIOT is free open-source, developed by a grassroots community

Comparison of Current Operating Systems								
OS	Min RAM	Min ROM	C Support	C++ Support	Multi-Threading	MCU w/o MMU	Modularity	Real-Time
Contiki	< 2kB	< 30kB	●	✗	●	✓	●	●
Tiny OS	< 1kB	< 4KB	✗	✗	●	✓	✗	✗
Linux	~ 1MB	~ 1MB	✓	✓	✓	✗	●	●
RIOT	~ 1.5kB	~ 5kB	✓	✓	✓	✓	✓	✓

Full support ✓

Partial support ●

No support ✗

# This Course – Two Parts

- Part One – Pr. Olivier Gruber – 50%
  - 5 weeks, widening your horizon as developers
  - Wandering outside of the Linux world
  - Also a primer before part two
- Part Two – Pr. Didier Donsez – 50%
  - 5 weeks, experimenting with an IoT infrastructure
  - Sensors → Gateway → Data-gathering Server



# First Part – Overview

- Widening your horizon as a developer
  - What you probably know about...



Application programming, right?

C, Java, JavaScript, Python...

Basic usage of your operating system, right?

File system and Graphical User Interface  
Maybe a hint of shell usage/scripting...

**How do we program these?**

**Which tools do we use?**





# First Part – Overview



## Bare-metal programming

Low-level C programming

Lucky if using an Hardware Abstraction Layer (HAL)

Otherwise needs some assembly language programming

## Embedded programming (real-time or specialized)

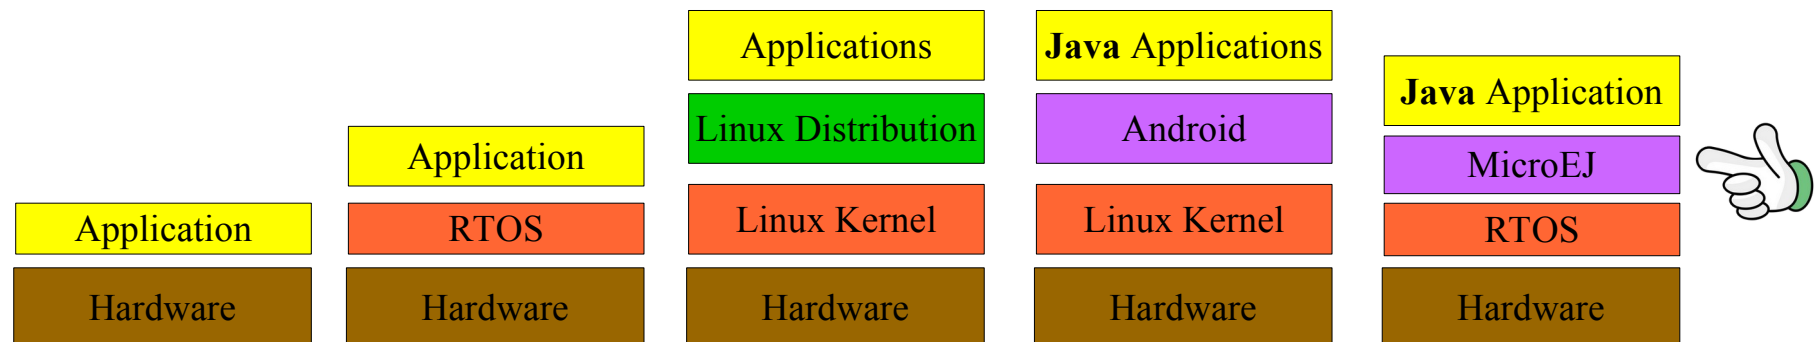
Higher-level programming... but not the usual Linux

Timing-aware programming...

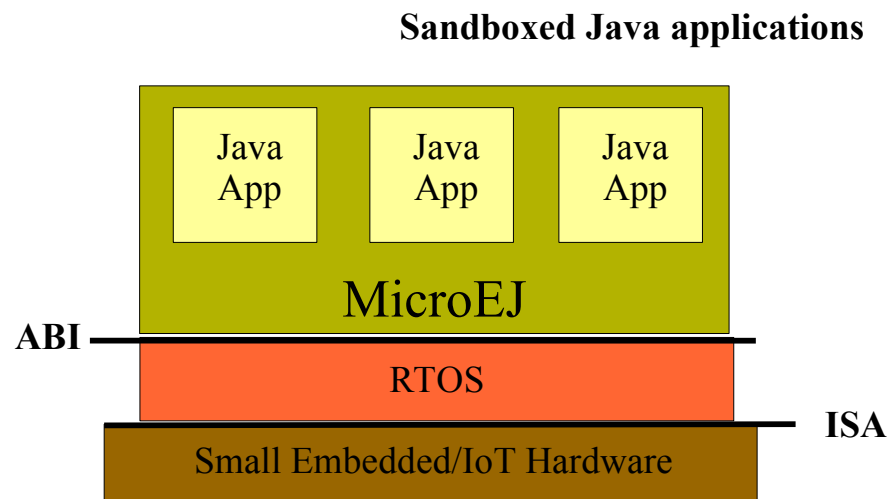
## Java programming

Increasingly more Java in the embedded world...

Not convinced? What about Android?



- Java for small and smart IoT devices
  - Similar virtualization technology as Android but *tighter* implementation
    - Android: \$15 processor, 32MB of memory
    - MicroEJ: \$1 processor, 128KB of memory
  - Google Cloud IoT Partners
    - IoT solutions that leverage Google's secure, global, and scalable infrastructure



# This Course Summary

## Part-I Bare-metal programming

Bare-metal programming

Making your own minimal Linux distribution

→ *Putting it all in perspective*



Applications

Mini-Distribution

Linux Kernel

Application

Hardware

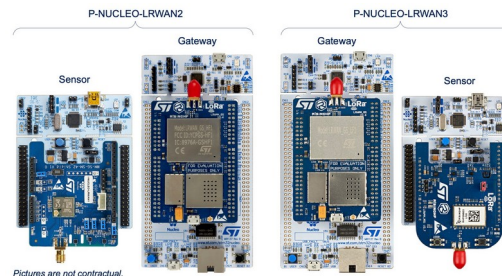
Hardware

## Part-II RTOS/IoT programming

RIOT programming

IoT Infrastructure

Lora networking



Application

RIOT

Hardware

# First Part – Pedagogy

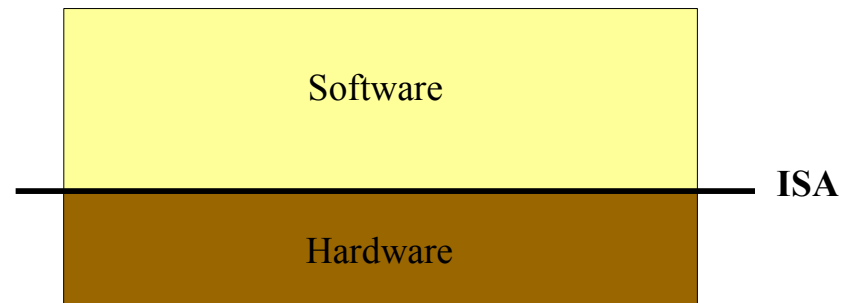
- Inverted pedagogy --- Team learning, individual work
  - Guided hands-on learning and coding
  - A work log to read, to understand, to complete, and to make your own
- The work log is for yourself first
  - A document about what you did and what you have learned
  - This document is first and foremost for your own records
- The destination is not the goal, the goal is the learning along that path
  - Don't just answer the questions...
  - Absorb the demonstrated know-how, ask questions about it, put it in practice...
  - Discuss with others what you have learned/understood
  - Help each others, the goal is the learning

# First Part – Evaluation

- Part-I Evaluation
  - No exam – but weekly progress checkpoints
  - Every week: you will surrender your work (work log and your code)
  - Your weekly involvement is the largest part of your final grade for this part of the M2M course

# Bare-metal Development

- Bare-metal Software
  - Runs directly on the "bare metal"
- Instruction Set Architecture (ISA)
  - Defines the instruction set
  - Defines other concepts such as page tables, traps, interrupts, etc.
- Cross-development Toolchain
  - Compiler, linker, debugger, and other tools



# Toolchain Background

- GNU toolchain

- GNU gcc, gdb, linker, and utility programs (binutils)
- Designed to work with Linux kernel

- Major hardware/software interfaces

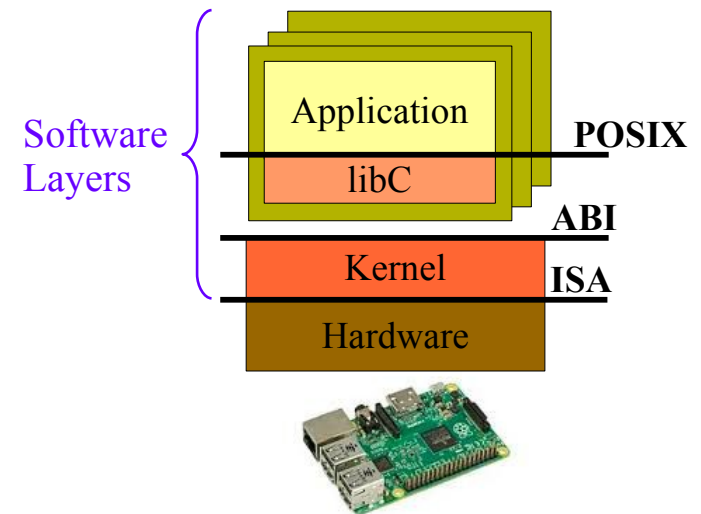
- ISA: Instruction Set Architecture
- (E)ABI: (*Embedded*) Application Binary Interface
- POSIX: Portable Operating System Interface standards

- Different GNU toolchains

- For different processors, different ABI, and different "kernels/distributions"
- Example: **i386-linux-abi** or **amd64-linux-abi**
- Each Linux distribution starts with a toolchain

- Used to build the kernel
- Used to build the libC
- Used to build the applications

} Different pieces of code, compiled independently, but interacting at runtime. Thus, the same toolchain must be used by all developers



# Cross-Compilation

- Cross compilation

- The development machine and the compilation target are different
- May have different ISA as well as different software stack



- Cross-compiling bare-metal software

- Developing on your AMD64 laptop, running Linux kernel, and an Ubuntu distribution
- Targetting a ARM board, like the VersatilePB or the Raspberry Pi
- Using different toolchains

- Raspi: gcc-8-arm-linux-gnueabi
- VersatilePB: **arm-none-eabi**

} Can be installed via *apt-get*<sup>(1)</sup>...  
If you need something else,  
you need to build your own toolchain...

(1) Using the program "synaptic" for a nice GUI over apt-get that allows you to search for available packages



# Bare-metal Development

- Bare-metal Software

- Runs directly on the "bare metal"

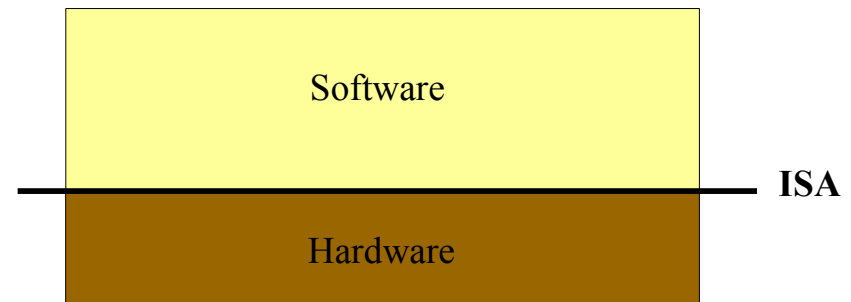
- Instruction Set Architecture (ISA)



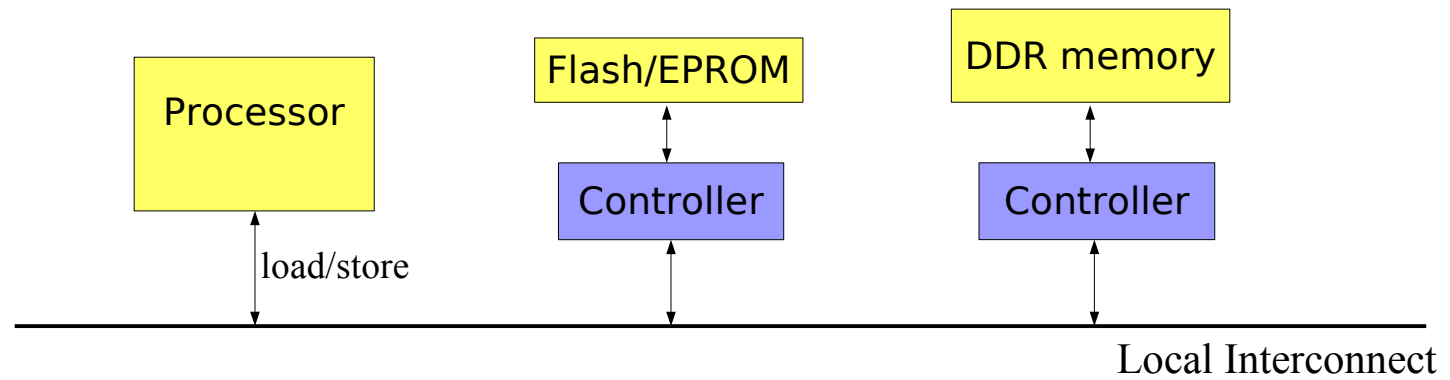
- Defines the instruction set
  - Defines other concepts such as page tables, traps, interrupts, etc.

- Cross-development Toolchain

- Compiler, linker, debugger, and other tools

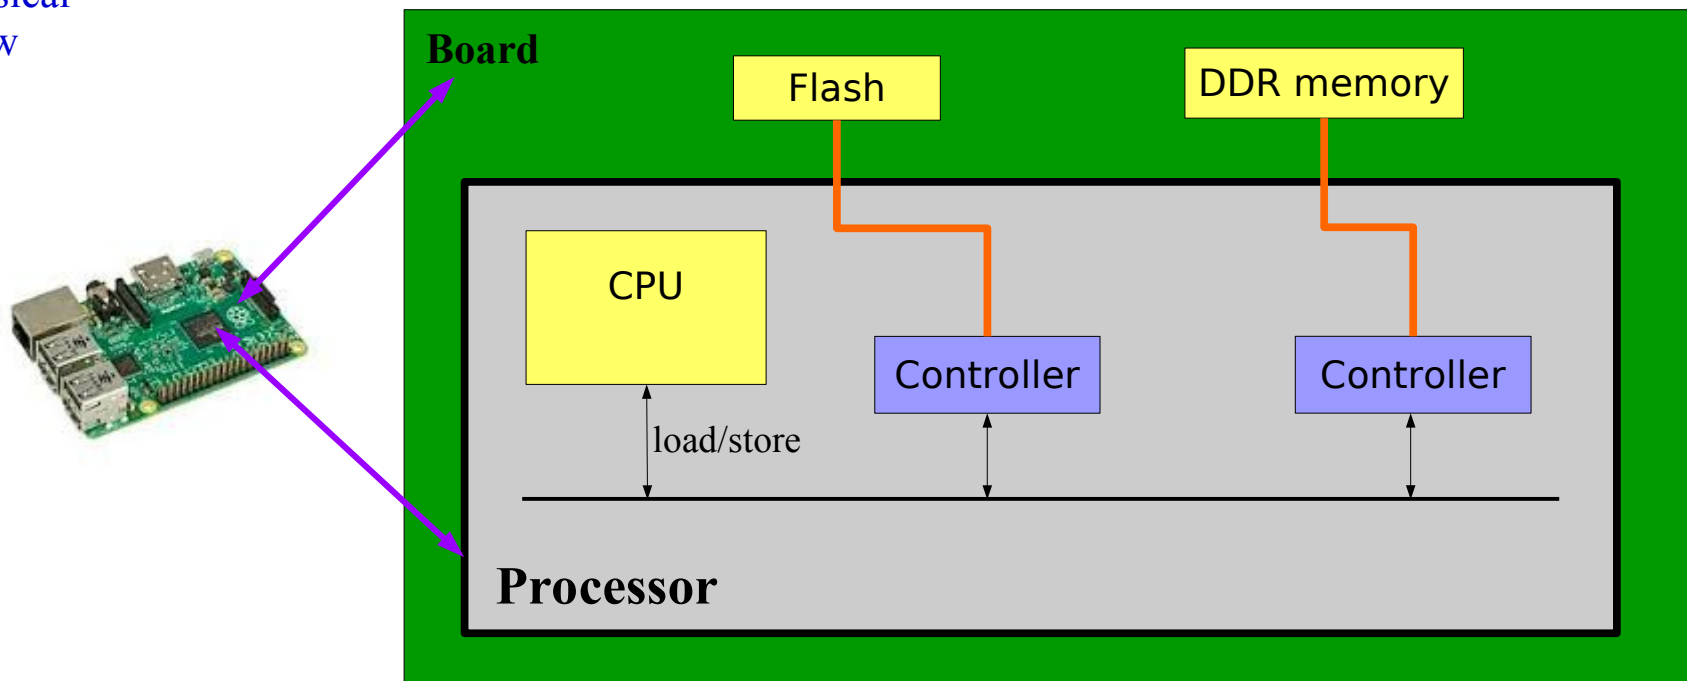


# Hardware – Basics

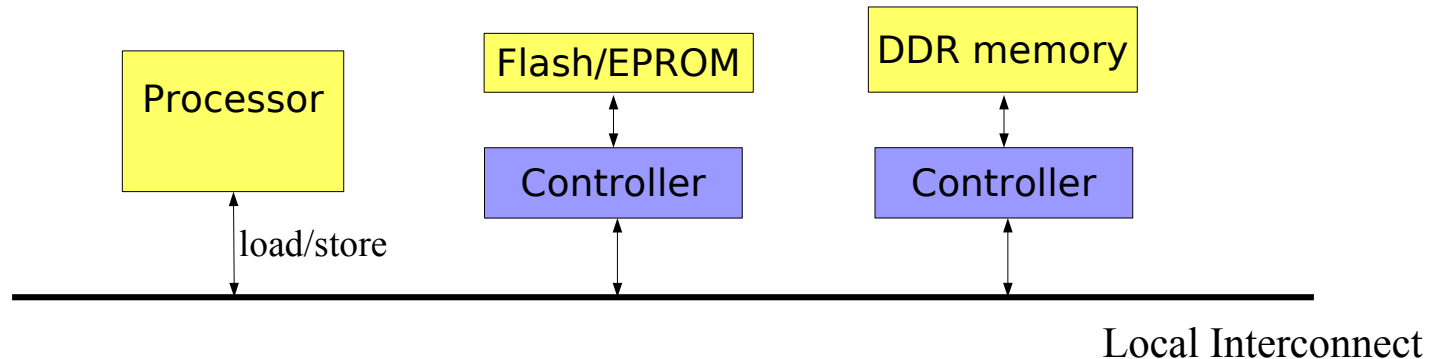


Conceptual  
View

Physical  
View

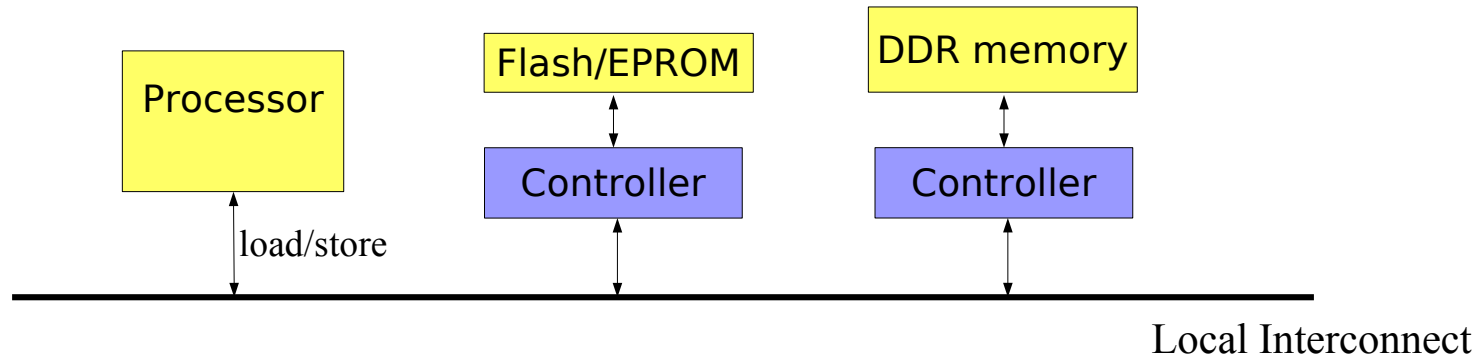


# Hardware – Reset / Power-up



- **Processor**
  - Only knows how to issue *load/store operations* on the bus
  - Forever loop: *fetch – decode – issue*
- **Local Interconnect (also called bus)**
  - Data wires + control wires
  - Routes load/store operations to controllers
  - Routing is based on the *memory map*

# Hardware – Reset / Power-up



- Boot sequence – ARM example
  - Processor wakes up at a given address, at 0x0000-0000 (reset vector)
  - Starts executing there, but what is there?

# ARM Example – Hardware/Software Boundary

21

- Page 0x0000-0000
  - Interrupts/Traps are *exceptions* for the processor

The 7 known exceptions,  
room for one 4-byte  
instructions per exception

```
.globl _vector = 0x00000000
_vector:
```

```
    ldr    pc, _reset
```

```
    ldr    pc, _undefined_instruction
```

```
    ldr    pc, _software_interrupt
```

```
    ldr    pc, _prefetch_abort
```

```
    ldr    pc, _data_abort
```

```
    ldr    pc, _not_used
```

```
    ldr    pc, _irq
```

```
    ldr    pc, _fiq
```

↔ ldr pc, [pc+18]

```
_reset:                                .word reset
```

```
_undefined_instruction: .word undefined_instruction_handler
```

```
_software_interrupt:   .word software_interrupt_handler
```

```
_prefetch_abort:      .word prefetch_abort_handler
```

```
_data_abort:          .word data_abort_handler
```

```
_not_used:            .word not_used_handler
```

```
_irq:                 .word irq_handler
```

```
_fiq:                 .word fiq
```

```
reset:
```

```
...
```

But what is the content of memory  
after powering up?

# Zynq-7000 Memory Map

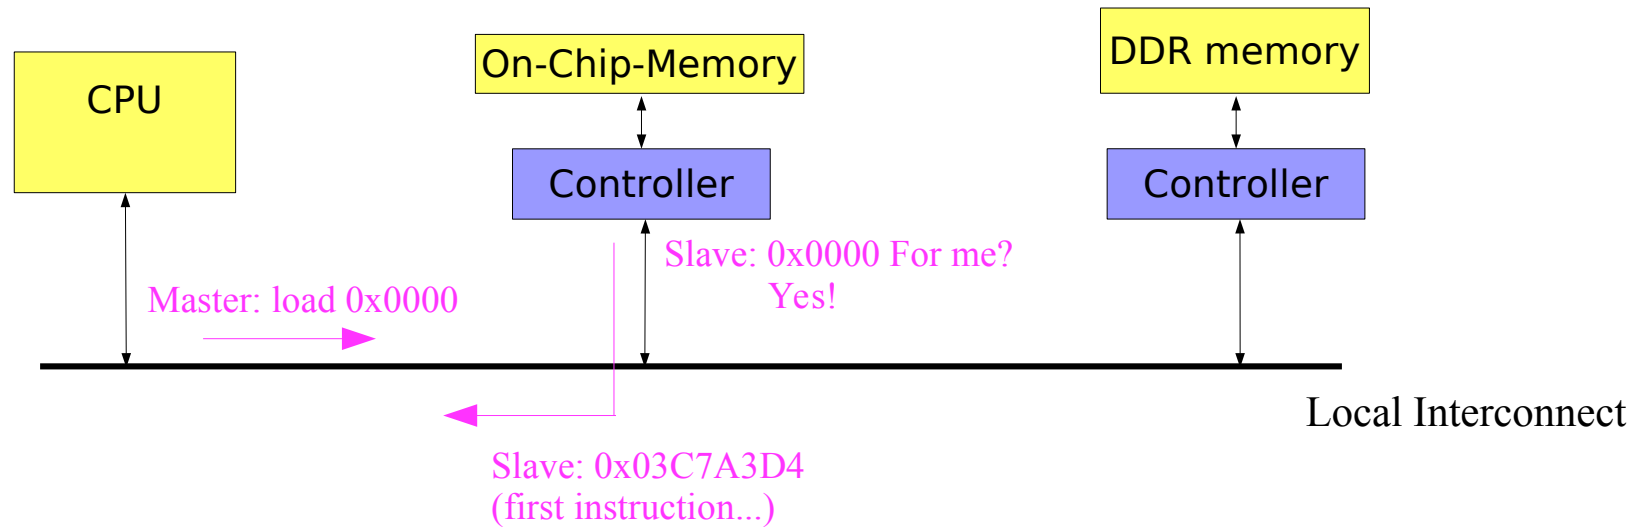
Table 4-1: System-Level Address Map

Address Range	CPUs and ACP	AXI_HP	Other Bus Masters <sup>(1)</sup>	Notes
0000_0000 to 0003_FFFF <sup>(2)</sup>	OCM	OCM	OCM	Address not filtered by SCU and OCM is mapped low
	DDR	OCM	OCM	Address filtered by SCU and OCM is mapped low
	DDR			Address filtered by SCU and OCM is not mapped low
				Address not filtered by SCU and OCM is not mapped low
0004_0000 to 0007_FFFF	DDR			Address filtered by SCU
				Address not filtered by SCU
0008_0000 to 000F_FFFF	DDR	DDR	DDR	Address filtered by SCU
		DDR	DDR	Address not filtered by SCU <sup>(3)</sup>
0010_0000 to 3FFF_FFFF	DDR	DDR	DDR	Accessible to all interconnect masters
4000_0000 to 7FFF_FFFF	PL		PL	General Purpose Port #0 to the PL, M_AXI_GP0
8000_0000 to BFFF_FFFF	PL		PL	General Purpose Port #1 to the PL, M_AXI_GP1
E000_0000 to E02F_FFFF	IOP		IOP	I/O Peripheral registers, see <a href="#">Table 4-6</a>
E100_0000 to E5FF_FFFF	SMC		SMC	SMC Memories, see <a href="#">Table 4-5</a>
F800_0000 to F800_0BFF	SLCR		SLCR	SLCR registers, see <a href="#">Table 4-3</a>
F800_1000 to F880_FFFF	PS		PS	PS System registers, see <a href="#">Table 4-7</a>
F890_0000 to F8F0_2FFF	CPU			CPU Private registers, see <a href="#">Table 4-4</a>
FC00_0000 to FDFF_FFFF <sup>(4)</sup>	Quad-SPI		Quad-SPI	Quad-SPI linear address for linear mode
FFFC_0000 to FFFF_FFFF <sup>(2)</sup>	OCM	OCM	OCM	OCM is mapped high
				OCM is not mapped high

Addresses are not filtered by the Snoop Control Unit (not configured yet)

On-Chip-Memory is mapped low, so the low addresses have meaningful contents

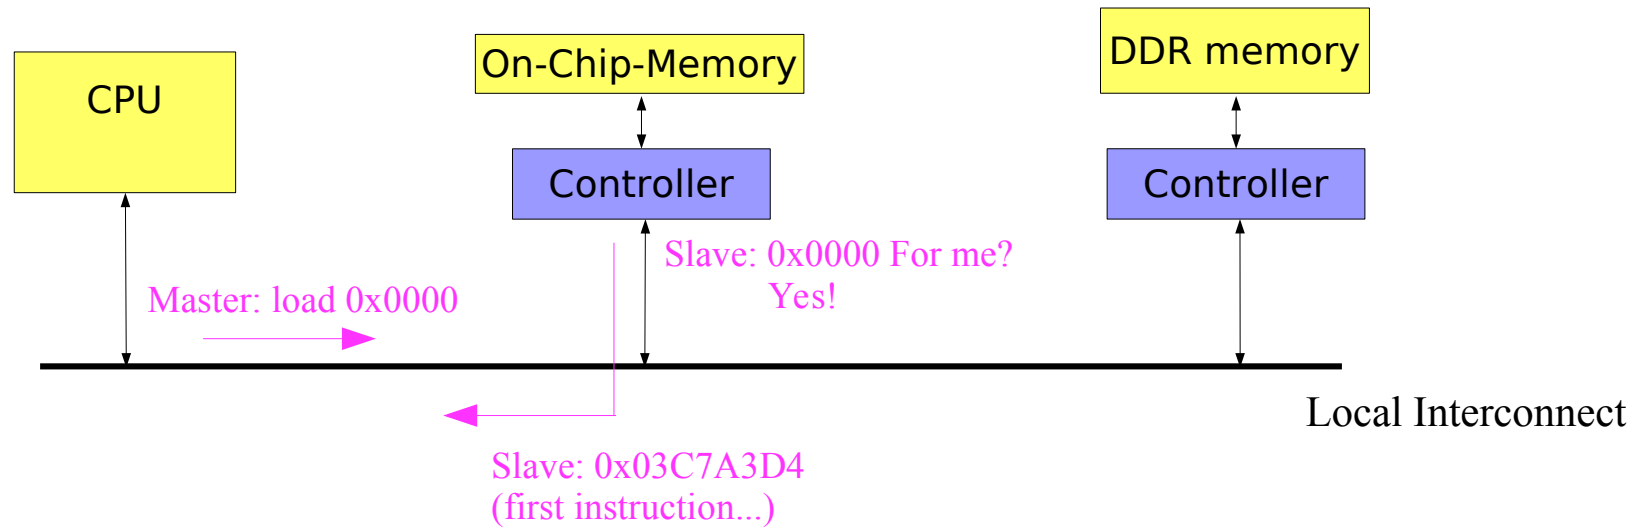
# Zybo – Reset / Power-up



- Primary boot sequence steps

- Reset the processor sets the program counter register to 0x0000-0000
- OCM is mapped low (0x0000-0000), so execute program stored in the OCM
- The code will initialize the various controllers/devices
- That code may be the only code that will ever run (like in a small IoT sensor)
- Or it may launch a secondary boot sequence, after having remapped the OCM high

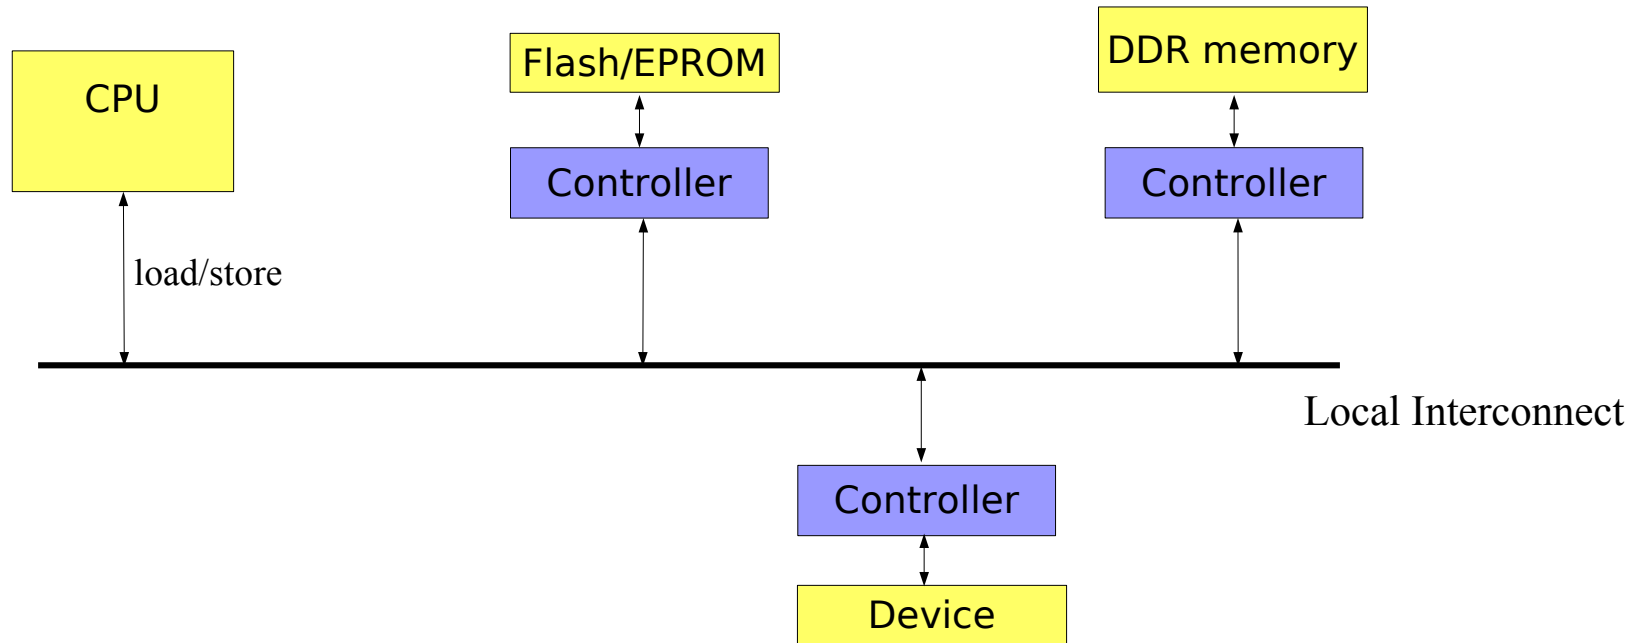
# Zybo – Reset / Power-up



- Secondary optional boot sequence steps
  - Loading some larger software from some mass-storage-like device, like an SD card
    - On Intel machine, the primary boot sequence is the BIOS
    - Then goes on loading a boot loader or a Linux kernel for example
  - Copies a trampoline in DDR and jumps to it
  - The trampoline allows to remap the OCM high and then jump to the loaded code



# Hardware – Questions

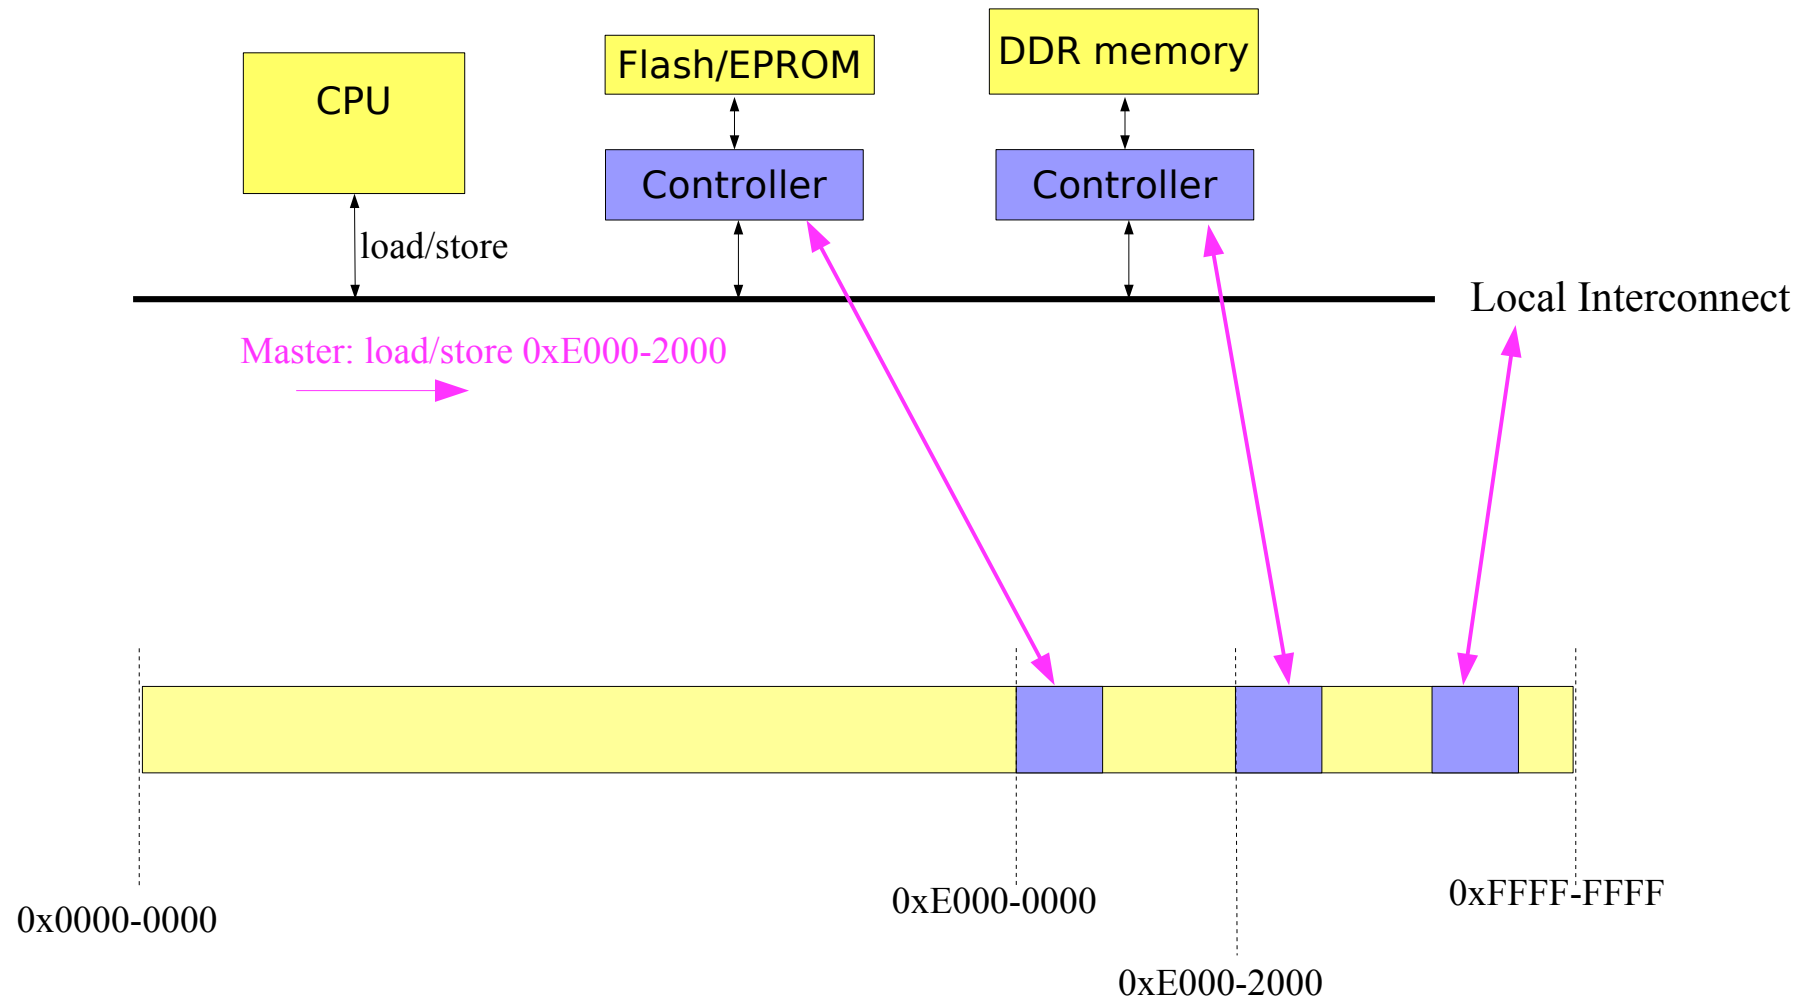


How can the Flash/EPROM can be remaped?

How can software configure hardware controllers?

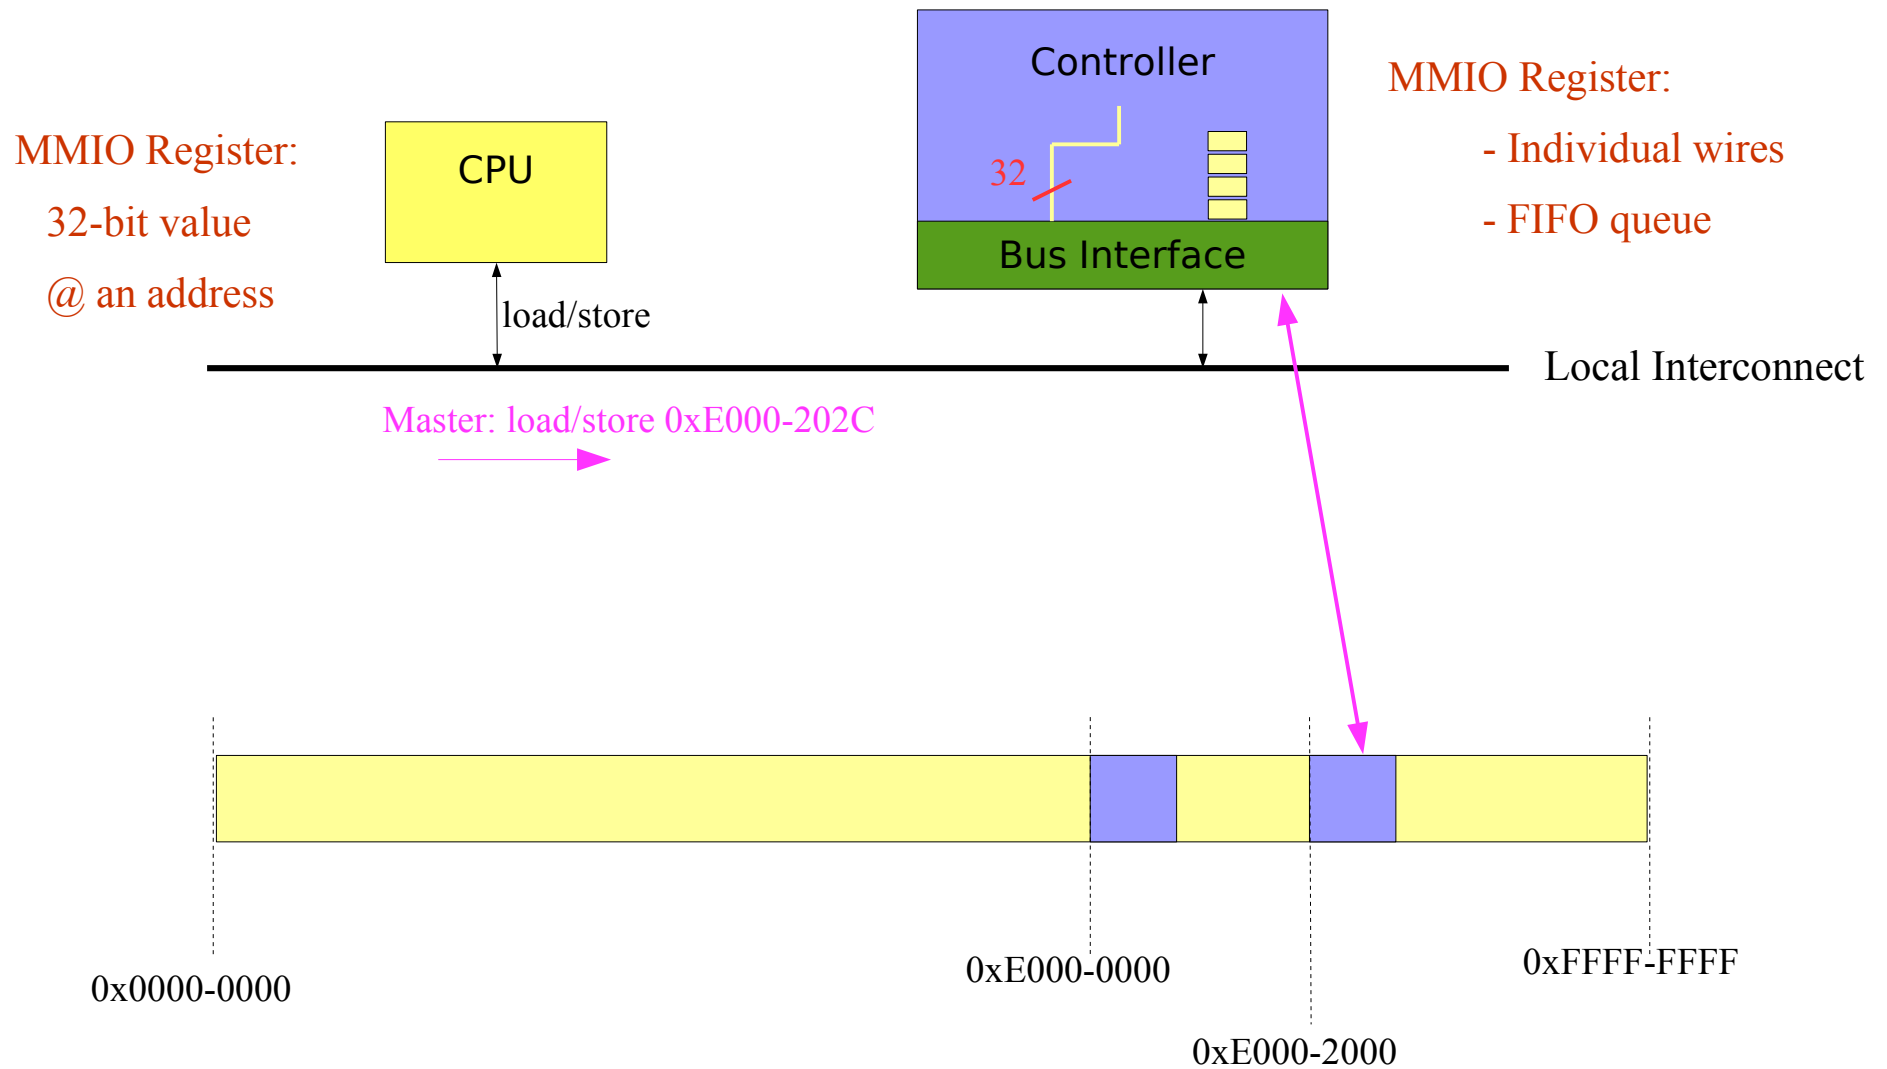
More generally, how does software interacts with devices attached to hardware controllers?

# Hardware – MMIO Registers



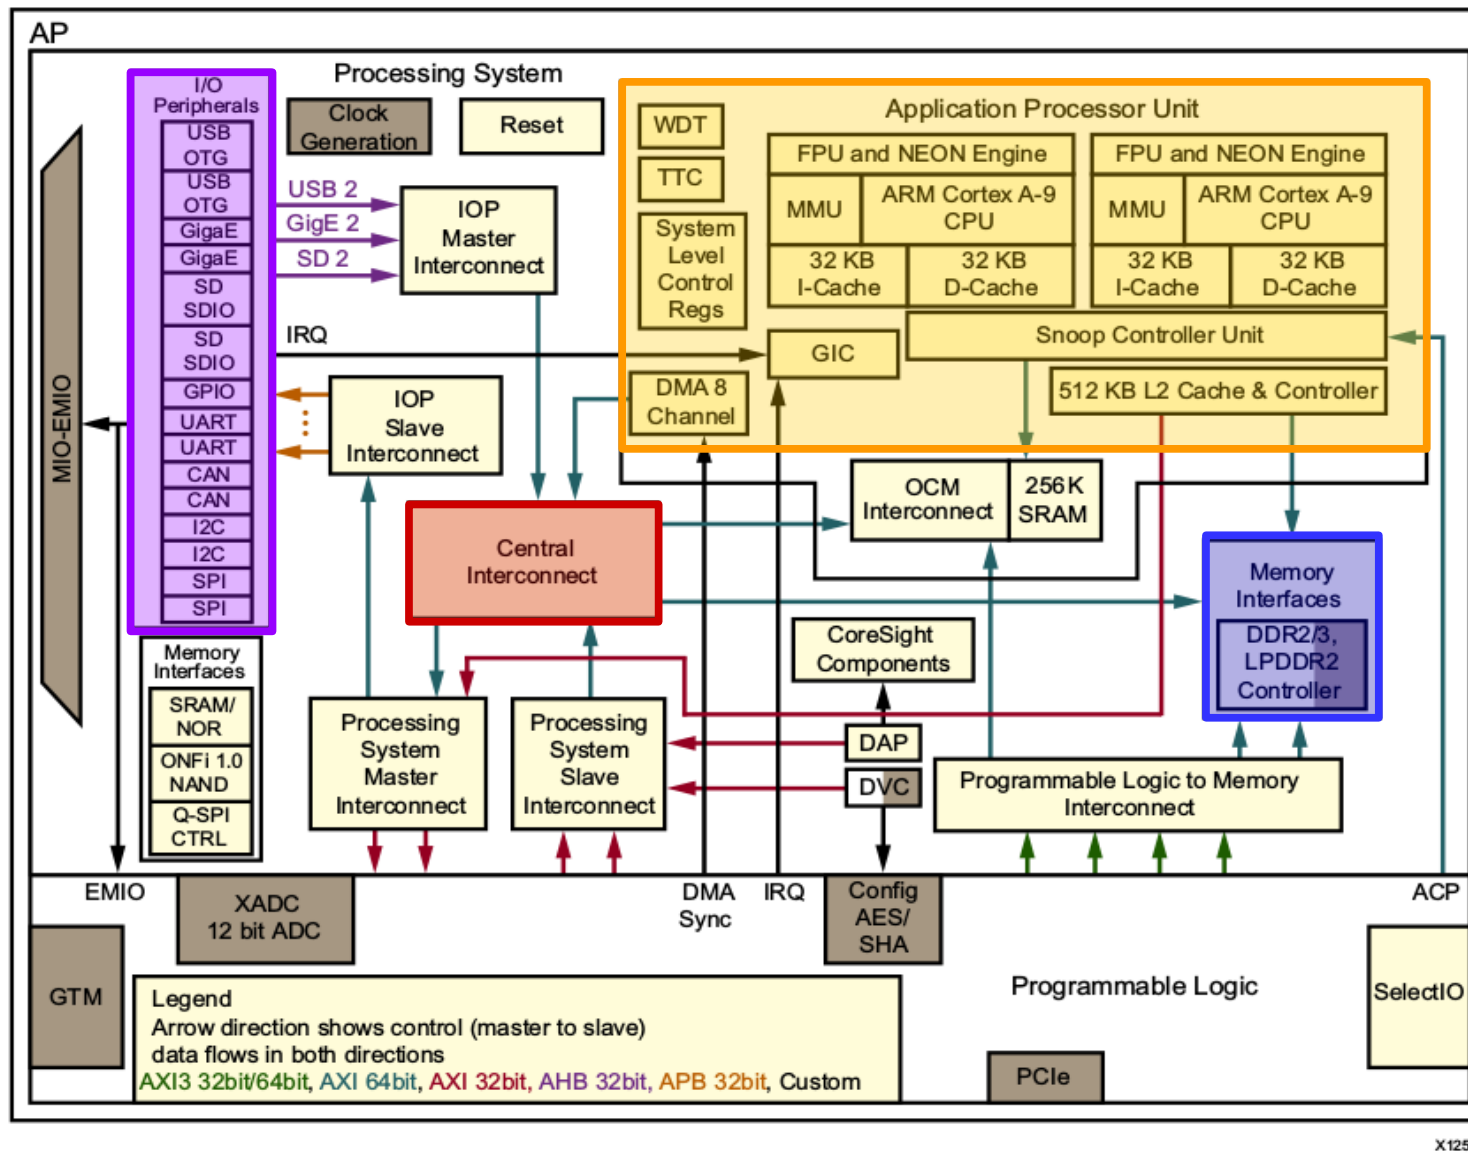
Memory-Mapped Input/Output Ranges

# Hardware – MMIO Registers



MMIO: Memory-Mapped Input/Output

# Hardware – Zynq-7000 Overview



# Zynq-7000 Memory Map

Table 4-1: System-Level Address Map

Address Range	CPUs and ACP	AXI_HP	Other Bus Masters <sup>(1)</sup>	Notes
0000_0000 to 0003_FFFF <sup>(2)</sup>	OCM	OCM	OCM	Address not filtered by SCU and OCM is mapped low
	DDR	OCM	OCM	Address filtered by SCU and OCM is mapped low
	DDR			Address filtered by SCU and OCM is not mapped low
				Address not filtered by SCU and OCM is not mapped low
0004_0000 to 0007_FFFF	DDR			Address filtered by SCU
				Address not filtered by SCU
0008_0000 to 000F_FFFF	DDR	DDR	DDR	Address filtered by SCU
		DDR	DDR	Address not filtered by SCU <sup>(3)</sup>
0010_0000 to 3FFF_FFFF	DDR	DDR	DDR	Accessible to all interconnect masters
4000_0000 to 7FFF_FFFF	PL		PL	General Purpose Port #0 to the PL, M_AXI_GP0
8000_0000 to BFFF_FFFF	PL		PL	General Purpose Port #1 to the PL, M_AXI_GP1
E000_0000 to E02F_FFFF	IOP		IOP	I/O Peripheral registers, see <a href="#">Table 4-6</a>
E100_0000 to E5FF_FFFF	SMC		SMC	SMC Memories, see <a href="#">Table 4-5</a>
F800_0000 to F800_0BFF	SLCR		SLCR	SLCR registers, see <a href="#">Table 4-3</a>
F800_1000 to F880_FFFF	PS		PS	PS System registers, see <a href="#">Table 4-7</a>
F890_0000 to F8F0_2FFF	CPU			CPU Private registers, see <a href="#">Table 4-4</a>
FC00_0000 to FDFF_FFFF <sup>(4)</sup>	Quad-SPI		Quad-SPI	Quad-SPI linear address for linear mode
FFFC_0000 to FFFF_FFFF <sup>(2)</sup>	OCM	OCM	OCM	OCM is mapped high
				OCM is not mapped high

# Zynq-7000 Memory Map

Table 4-7: PS System Register Map

Register Base Address	Description (Acronym)	Register Set
F800_1000, F800_2000	Triple timer counter 0, 1 (TTC 0, TTC 1)	ttc.
F800_3000	DMAC when secure (DMAC S)	dmac.
F800_4000	DMAC when non-secure (DMAC NS)	dmac.
F800_5000	System watchdog timer (SWDT)	swdt.
F800_6000	DDR memory controller	ddrc.
F800_7000	Device configuration interface (DevC)	devcfg.
F800_8000	AXI_HP 0 high performance AXI interface w/ FIFO	afi.
F800_9000	AXI_HP 1 high performance AXI interface w/ FIFO	afi.
F800_A000	AXI_HP 2 high performance AXI interface w/ FIFO	afi.
F800_B000	AXI_HP 3 high performance AXI interface w/ FIFO	afi.
F800_C000	On-chip memory (OCM)	ocm.
F800_D000	eFuse <sup>(1)</sup>	-
F800_F000	Reserved	Reserved

# Zynq-7000 Memory Map

*Table 4-6: I/O Peripheral Register Map*

Register Base Address	Description
E000_0000, E000_1000	UART Controllers 0, 1
E000_2000, E000_3000	USB Controllers 0, 1
E000_4000, E000_5000	I2C Controllers 0, 1
E000_6000, E000_7000	SPI Controllers 0, 1
E000_8000, E000_9000	CAN Controllers 0, 1
E000_A000	GPIO Controller
E000_B000, E000_C000	Ethernet Controllers 0, 1
E000_D000	Quad-SPI Controller
E000_E000	Static Memory Controller (SMC)
E010_0000, E010_1000	SDIO Controllers 0, 1

**UART** = Universal Asynchronous Receiver and Transmitter

**Also called RS-232** (a standard for a serial line over 2 wires)

# UART Device Example

**UART... serial line controller, following the RS-232 protocol...**

**→ Essentially a FIFO and a status register...**

Corresponding the **mmio** registers defined in the Zynq-7000/R1P8 Technical Reference Manual

```
#define UART_R1P8_CR          0x0000 /* UART Control Register */
#define UART_R1P8_MR          0x0004 /* UART Mode Register */
#define UART_R1P8_IER         0x0008 /* -- Interrupt Enable Register */
#define UART_R1P8_IDR         0x000C /* -- Interrupt Disable Register */
#define UART_R1P8_IMR         0x0010 /* -- Interrupt Mask Register */
#define UART_R1P8_ISR         0x0014 /* -- Channel Interrupt Status Register */
#define UART_R1P8_BAUDGEN     0x0018 /* Baude Rate Generator Register */
#define UART_R1P8_RXTOUT      0x001C /* -- Receiver Timeout Register */
#define UART_R1P8_RXWM        0x0020 /* -- Receiver FIFO Trigger Level Register */
#define UART_R1P8_MODEMCR     0x0024 /* -- Modem Control Register */
#define UART_R1P8_MODEMSR     0x0028 /* -- Modem Status Register */
#define UART_R1P8_SR          0x002C /* Channel Status Register */
#define UART_R1P8_FIFO        0x0030 /* Transmit & Receive FIFO */
#define UART_R1P8_BAUDDIV     0x0034 /* Baud Rate Divider Register */
#define UART_R1P8_FLOWD       0x0038 /* -- Flow Control Delay Register */
#define UART_R1P8_TXWM        0x0044 /* -- Transmitter FIFO Trigger Level Register */
```



# UART Device Example

## Interacting with a device:

- 1) choose one mmio range corresponding to your device
- 2) choose one or more register (at different offsets in that range)
- 3) work with one or more bits in that register

```
#define UART0 0xE0000000
#define UART1 0xE0001000
```

```
#define UART_R1P8_SR          0x002C /* Channel Status Register */
#define UART_R1P8_FIFO        0x0030 /* Transmit & Receive FIFO */
/*
```

```
 * Channel Status Register (UART_R1P8_SR)
 */
```

```
#define UART_R1P8_SR_TNFUL  (1 << 14)
#define UART_R1P8_SR_TTRIG  (1 << 13)
#define UART_R1P8_SR_FDELT  (1 << 12)
#define UART_R1P8_SR_TACTIVE (1 << 11)
#define UART_R1P8_SR_RACTIVE (1 << 10)
```

```
#define UART_R1P8_SR_TFUL    (1 << 4)
#define UART_R1P8_SR_EMPTY  (1 << 3)
#define UART_R1P8_SR_RFUL   (1 << 2)
#define UART_R1P8_SR_REMPTY (1 << 1)
#define UART_R1P8_SR_RTRIG  (1 << 0)
```

# UART – Initialization

```
void
uart_r1p8_init_regs(void* uart) { /* See Zynq TRM sequence (UG585 p598) */

    /* UART Character frame */
    mmio_reg_write32(uart,UART_R1P8_MR,UART_R1P8_MR_8n1);

    /* Baud Rate configuration */
    mmio_reg_setbits32(uart,UART_R1P8_CR,  UART_R1P8_CR_RXDIS | UART_R1P8_CR_TXDIS);
    mmio_reg_write32(uart,UART_R1P8_BAUDGEN, UART_R1P8_115200_GEN);
    mmio_reg_write32(uart,UART_R1P8_BAUDDIV, UART_R1P8_115200_DIV);
    mmio_reg_setbits32(uart,UART_R1P8_CR,  UART_R1P8_CR_RXRES | UART_R1P8_CR_TXRES);
    mmio_reg_setbits32(uart,UART_R1P8_CR,  UART_R1P8_CR_RXEN | UART_R1P8_CR_TXEN);

    /* Disable Rx Trigger level */
    mmio_reg_write32(uart,UART_R1P8_RXWM,  0x00);

    /* Enable Controller */
    mmio_reg_write32(uart,UART_R1P8_CR,  UART_R1P8_CR_RXRES | UART_R1P8_CR_TXRES |
        UART_R1P8_CR_RSTTO | UART_R1P8_CR_RXEN | UART_R1P8_CR_TXEN |
        UART_R1P8_CR_STPBRK);

    /* Configure Rx Timeout */
    mmio_reg_write32(uart,UART_R1P8_RXTOUT, 0x00);
    mmio_reg_write32(uart,UART_R1P8_IER, 0x00);
    mmio_reg_write32(uart,UART_R1P8_IDR,  UART_R1P8_IxR_ALL);

    /* No Flow delay */
    mmio_reg_write32(uart,UART_R1P8_FLOWD, 0x00);

    /* Deactivate flowcontrol */
    mmio_reg_clearbits32(uart,UART_R1P8_MODEMCR, UART_R1P8_MODEMCR_FCM);

    /* Mask all interrupts */
    mmio_reg_clearbits32(uart,UART_R1P8_IMR, 0x01FFF);
}
```

# UART – Output

```
#define UART0 0xE0000000
#define UART1 0xE0001000
```

```
#define UART_R1P8_SR      0x002C /* Channel Status Register */
#define UART_R1P8_FIFO    0x0030 /* Transmit & Receive FIFO */
```

```
/*
 * Channel Status Register (UART_R1P8_SR)
 */
#define UART_R1P8_SR_TNFUL  (1 << 14)
#define UART_R1P8_SR_TTRIG  (1 << 13)
#define UART_R1P8_SR_FDELTA (1 << 12)
#define UART_R1P8_SR_TACTIVE (1 << 11)
#define UART_R1P8_SR_RACTIVE (1 << 10)
#define UART_R1P8_SR_TFUL   (1 << 4)
#define UART_R1P8_SR_EMPTY  (1 << 3)
#define UART_R1P8_SR_RFUL   (1 << 2)
#define UART_R1P8_SR_REMPTY  (1 << 1)
#define UART_R1P8_SR_RTRIG   (1 << 0)
```

```
void
uart_r1p8_putc(void* uart, uint8_t c) {
    while((mmio_read32(uart, UART_R1P8_SR) & UART_R1P8_SR_TFUL) != 0)
        ;
    mmio_write32(uart, UART_R1P8_FIFO, c);
}
```

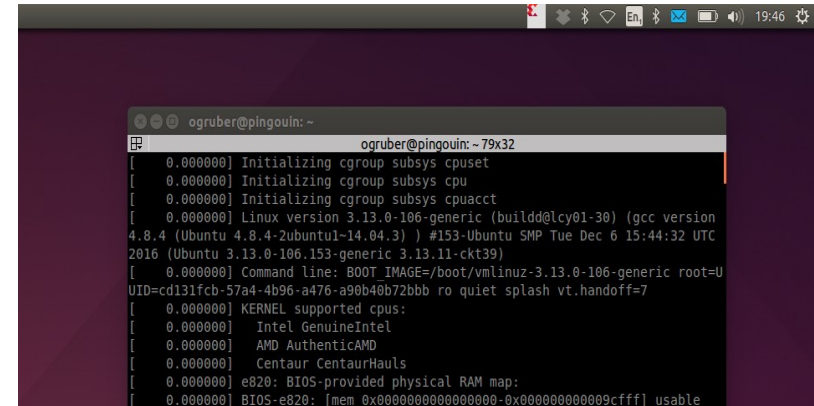
# UART – Input

```
#define UART_R1P8_SR          0x002C /* Channel Status Register */
#define UART_R1P8_FIFO        0x0030 /* Transmit & Receive FIFO */
/*
 * Channel Status Register (UART_R1P8_SR)
 */
#define UART_R1P8_SR_TNFUL    (1 << 14)
#define UART_R1P8_SR_TTRIG    (1 << 13)
#define UART_R1P8_SR_FDELT    (1 << 12)
#define UART_R1P8_SR_TACTIVE  (1 << 11)
#define UART_R1P8_SR_RACTIVE  (1 << 10)
#define UART_R1P8_SR_TFUL     (1 << 4)
#define UART_R1P8_SR_EMPTY    (1 << 3)
#define UART_R1P8_SR_RFUL     (1 << 2)
#define UART_R1P8_SR_EMPTY    (1 << 1)
#define UART_R1P8_SR_RTRIG    (1 << 0)
```

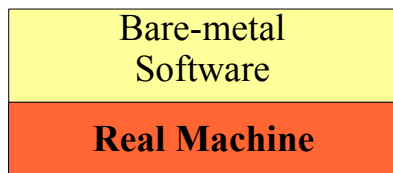
```
uint8_t
uart_r1p8_getc(void* uart){
    while((mmio_read32(uart, UART_R1P8_SR) & UART_R1P8_SR_EMPTY))
        ;
    return mmio_read32(uart, UART_R1P8_FIFO);
}
```

# Development Environment

- Using a real board
  - Relies on using a USB-Serial cable
    - Shows as `/dev/ttyUSB0` or `/dev/ttyUSB1`
    - Look at ***dmesg*** when plugging the cable in
  - Serial line used for the ***console***
    - The *standard input/output streams*
    - Requires using a terminal emulator on your laptop
    - Like ***minicom*** for instance
    - Typical configuration<sup>(1)</sup>: "115200 8N1"



```
ogrub@pinguin: ~  
ogrub@pinguin: ~ 79x32  
[ 0.000000] Initializing cgroup subsys cpuset  
[ 0.000000] Initializing cgroup subsys cpu  
[ 0.000000] Initializing cgroup subsys cpuacct  
[ 0.000000] Linux version 3.13.0-106-generic (build@lcy01-30) (gcc version 4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC 2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39)  
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=UUID=cd131fcb-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7  
[ 0.000000] KERNEL supported cpus:  
[ 0.000000] Intel GenuineIntel  
[ 0.000000] AMD AuthenticAMD  
[ 0.000000] Centaur CentaurHauls  
[ 0.000000] e820: BIOS-provided physical RAM map:  
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x000000000009cfff] usable
```

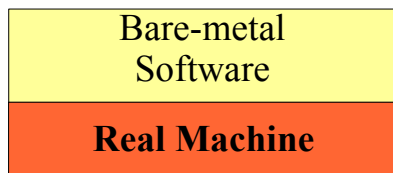


USB – Serial Cable – RS 232

(1) "115200 8N1": the baud rate is set to 115200bps, data-bits set to 8 (the '8' in 8N1), parity is set to none (the 'N' in 8N1), and stop-bits is set to 1 (the '1' in 8N1)

# Development Environment

- Hardware debugging
  - Through JTAG, often behind a USB connector
- Software debugging
  - Open On-Chip Debugger ([www.openocd.org](http://www.openocd.org))
  - Acts as a GDB debug server
  - But requires a *proper configuration file*<sup>(1)</sup>
  - Then can use gdb with a target remote connection
  - Can even load firmware/software directly at reset<sup>(2)</sup>



**USB Cable  
JTAG and Serial line**

A screenshot of a terminal window titled "ogruher@pinguin: ~". The terminal displays boot logs for a Linux system (Ubuntu 4.8.4-2ubuntu1-14.04.3) running on a Raspberry Pi. The logs show the initialization of cgroup subsystems, the kernel version (3.13.0-106-generic), and the command line used for booting. The terminal output includes: [ 0.000000] Initializing cgroup subsys cpuset, [ 0.000000] Initializing cgroup subsys cpu, [ 0.000000] Initializing cgroup subsys cpuacct, [ 0.000000] Linux version 3.13.0-106-generic (build@lcy01-30) (gcc version 4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC 2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39), [ 0.000000] Command line: BOOT\_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=UUID=c4131fcb-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7, [ 0.000000] KERNEL supported cpus: [ 0.000000] Intel GenuineIntel, [ 0.000000] AMD AuthenticAMD, [ 0.000000] Centaur CentaurHauls, [ 0.000000] e820: BIOS-provided physical RAM map: [ 0.000000] BIOS-e820: [mem 0x00000000-0x00000000-0x00000000-0x00000000] usable.



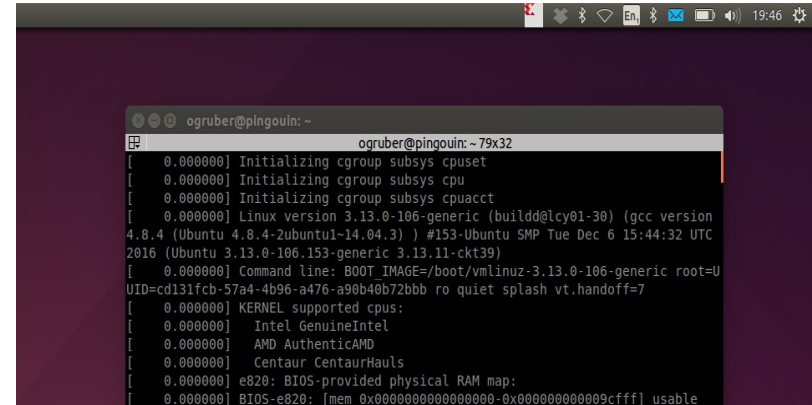
(1) This is the hard part, those are hard to find and not always correct  
(2) Avoids having to switching SD cards → will ultimately break the SD socket

# Development Environment

- Using an emulator – QEMU
  - A regular Linux process
  - Emulates a machine for your bare-metal software
    - A processor (intel or arm) with memory
    - A serial line for the console
    - One or more mass-storage devices for file systems
  - Support easy debugging
    - Provides a GDB debug server for GDB debugging

```
$ qemu-system-arm -serial mon:stdio -kernel your-code
```

```
$ qemu-system-i386 -hda disk.img
```



```
ogruher@pinguin: ~  
ogruher@pinguin: ~ 79x32  
[ 0.000000] Initializing cgroup subsys cpuset  
[ 0.000000] Initializing cgroup subsys cpu  
[ 0.000000] Initializing cgroup subsys cpuacct  
[ 0.000000] Linux version 3.13.0-106-generic (build@lcy01-30) (gcc version  
4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC  
2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39)  
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=U  
UID=cdf131fcb-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7  
[ 0.000000] KERNEL supported cpus:  
[ 0.000000] Intel GenuineIntel  
[ 0.000000] AMD AuthenticAMD  
[ 0.000000] Centaur CentaurHauls  
[ 0.000000] e820: BIOS-provided physical RAM map:  
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x0000000000009cfff] usable
```

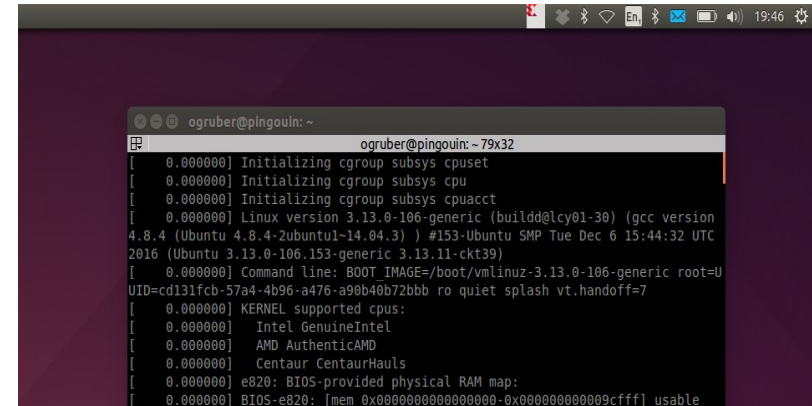


~~USB – Serial Cable – RS 232~~



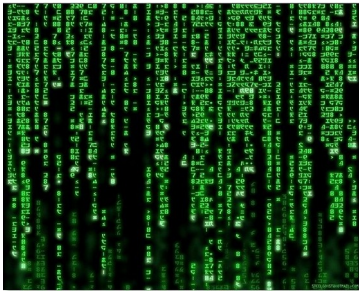
# Development Environment

- Using an emulator – QEMU
  - A regular Linux process
  - Emulates a machine for your bare-metal software
  - *But how is possible?*



```
ogruher@pinguin: ~  
ogruher@pinguin: ~ 79x32  
[ 0.000000] Initializing cgroup subsys cpuset  
[ 0.000000] Initializing cgroup subsys cpu  
[ 0.000000] Initializing cgroup subsys cpuacct  
[ 0.000000] Linux version 3.13.0-106-generic (build@lcy01-30) (gcc version  
4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC  
2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39)  
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=U  
UID=cd131fcb-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7  
[ 0.000000] KERNEL supported cpus:  
[ 0.000000] Intel GenuineIntel  
[ 0.000000] AMD AuthenticAMD  
[ 0.000000] Centaur CentaurHauls  
[ 0.000000] e820: BIOS-provided physical RAM map:  
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x00000000000009cfff] usable
```

**For your code:** "The matrix is a prison that you cannot see, taste, or smell." Morpheus



*For you: an all-software-development experience in the confort of your chair!*

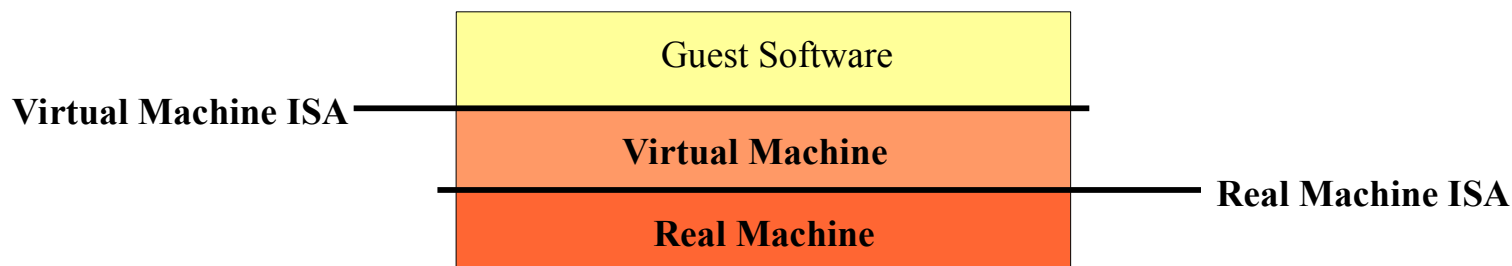


**USB – Serial Cable – RS 232**



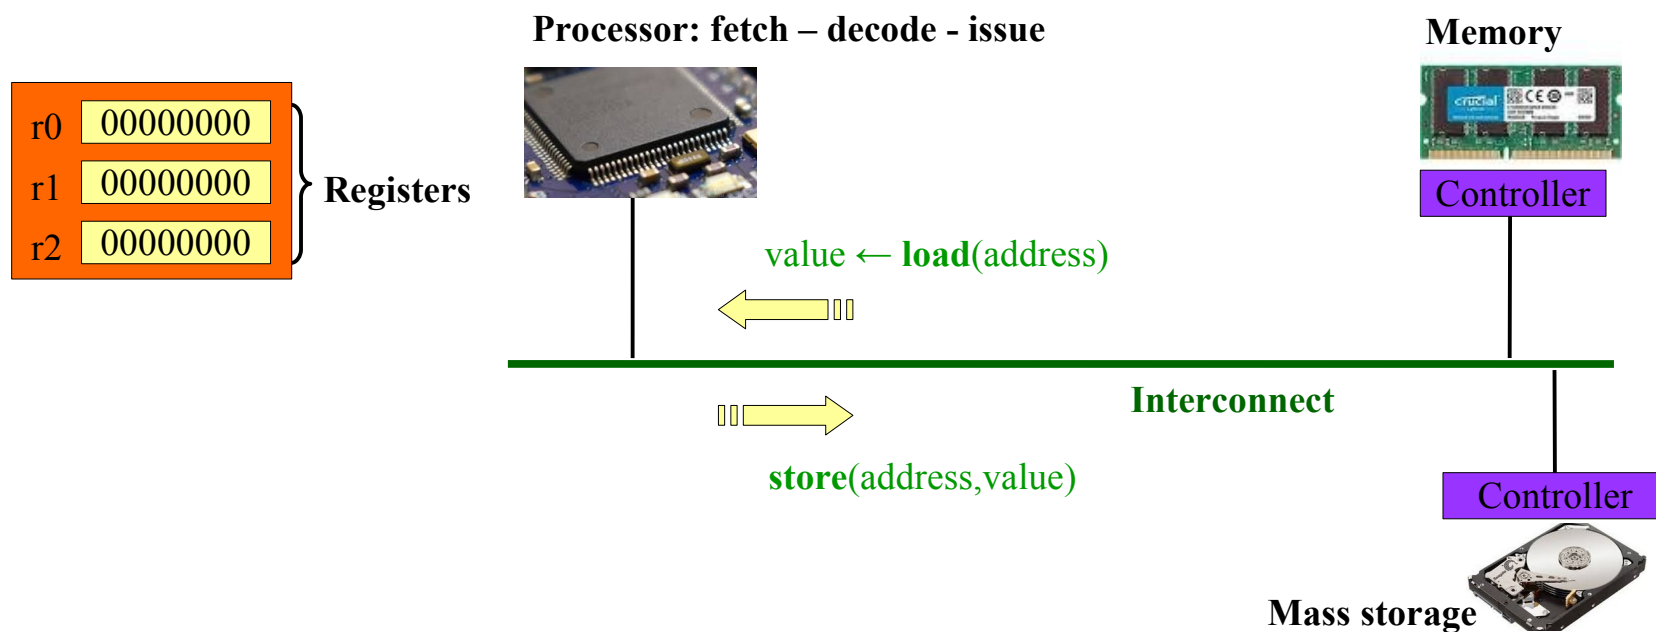
# Virtual Machines – The Art of Illusions

41



*What is an execution?*

*Imagine yourself being the code running on the processor, what can you see?*



# Illusion by Emulation

42

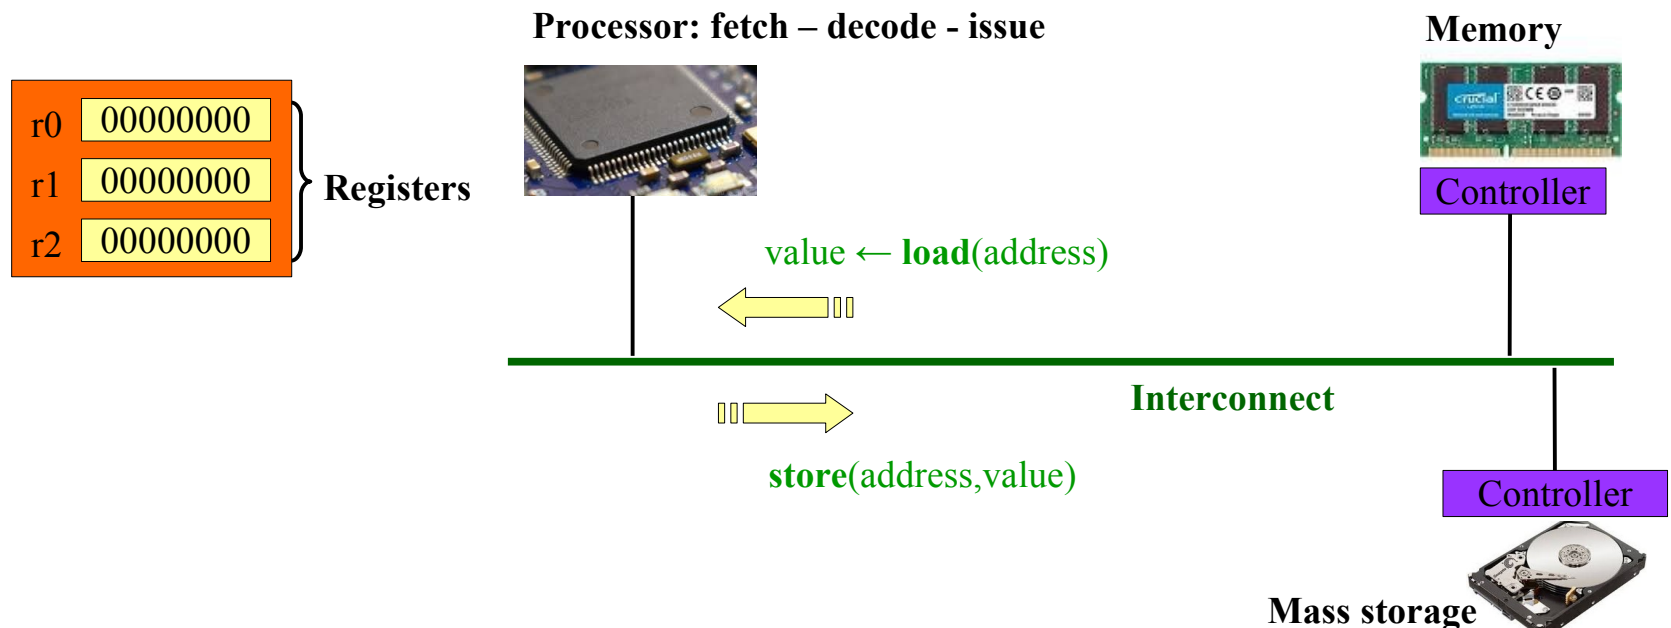
## *Any execution can be emulated by program*

- Fetch instructions, decode them, and issue them
- Maintain the contents of registers
- Maintain the contents of memory

(let's ignore devices for now...)

*After all, a **processor** is only an **interpreter** of instructions, done in **hardware**...*

*→ So it can be done in software.*



- **Architected State**

- Memory

- A number of physical frames, can be emulated by an array

- Processor – per core

- A set of registers (r0-r15 for example)
    - Current status: user or kernel mode
    - Interrupt status

- **Core Emulation**

- Just a regular program...

- Emulating the execution...

```
uint8_t memory[SIZE_512MB];
```

```
struct core {  
    int32_t regs[16];  
    uint32_t flags;  
    uint32_t pending_irqs;  
    struct mmu *mmu;  
};
```

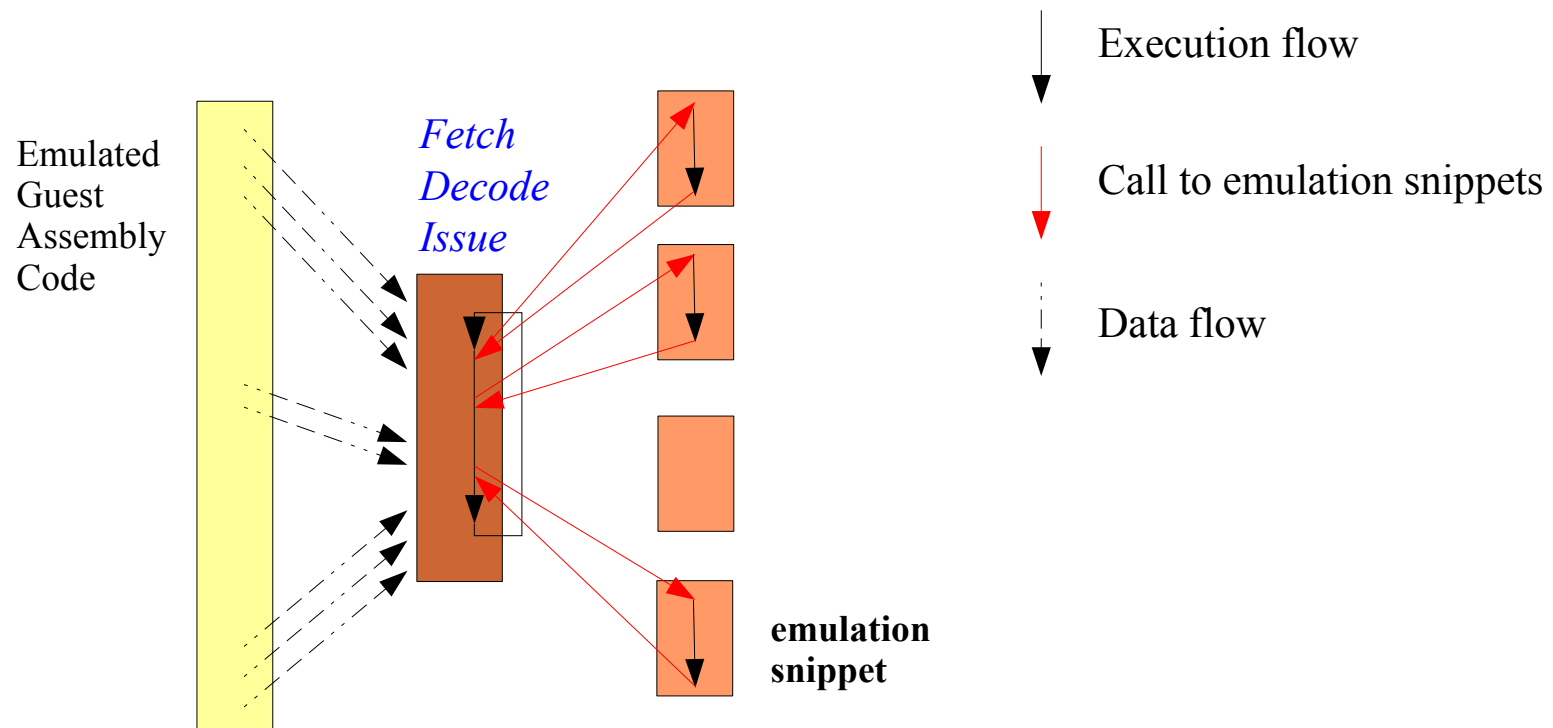
```
struct processor {  
    struct core cores[4];  
    ...  
};
```

```
void wakeup(struct processor *p, int coreno) {  
    reset();  
    for (;;) {  
        fetch_decode_issue();  
    }  
}
```

# Emulation by Interpretation

44

- The emulation is done via an interpreter
  - Like any other interpreter (Shell, JavaScript, or Java)
  - Emulate the Fetch-Decode-Issue loop (FDI loop)
    - Fetch instruction at emulated program counter in emulated memory
    - Decode the instruction → invoke emulation snippet
    - Emulation snippets manipulate the architected state



# Emulation Snippets – Examples

45

**mov r1,r0**

```
→ core.regs[dst] = core.regs[src];    // actual copy of the emulate register  
   core.regs[pc] += 4;                  // update the program counter
```

**add r1,r1,r0**

```
→ core.regs[dst] = core.regs[left] + core.regs[right]; // the arithmetic operation  
   core.regs[pc] += 4;                                     // update the program counter
```

**ldr r1,[r0]**

```
→ uint32_t* addr = core.regs[src];    // guest address to load from emulated memory  
   core.regs[dst] = memory[addr];      // emulate load from memory  
   core.regs[pc] += 4;                  // update the program counter
```

**bl r1**

```
→ uint32_t* addr = core.regs[src];    // guest address to branch to  
   core.regs[link] = core.regs[pc];   // save current program counter in link register  
   core.regs[pc] = *addr;               // branch the emulated execution
```

```
struct core {  
    int32_t regs[15];  
} core;
```

```
uint8_t memory[SIZE_64MB];
```

## Set the page table directory (turn on MMU)

```
mov mmu,r0      → core.ptd = core.regs[src];    // actual copy of the emulated register
                  core.regs[pc] += 4;           // update the program counter
```

## Memory load/store operations must emulate the MMU now

```
ldr r1,[r0]     → uint32_t* addr = core.regs[src];    // address to load from emulated memory
                  addr = softMMU_translate(addr);      // translate address from virtual to physical
                  core.regs[dst] = memory[addr];       // emulate load from memory
                  core.regs[pc] += 4;                 // update the program counter
```

```
struct page_table_entry {
    int  flags:8;
    int  addr:24;
};
```

```
struct page_table {
    struct page_table_entry* entries;
    int depth;
    int page_size;
};
```

```
struct core {
    int32_t regs[15];
    struct page_table *pdt;
} core;
```

```
uint8_t memory[SIZE_64MB];
```

# Virtual Machines – Type I or II

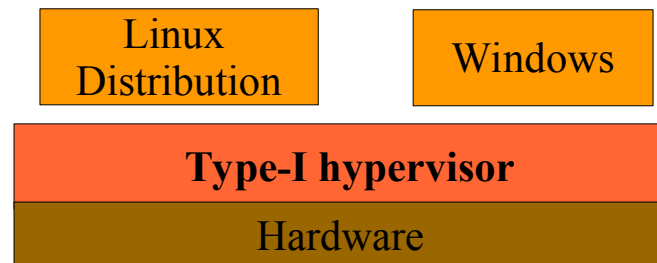
47

## **Type-I Hypervisors<sup>(1)</sup>:**

*VMware, Xen, ...*

*Installed directly on the hardware.*

*Host multiple guest operating systems that are para-virtualized.*



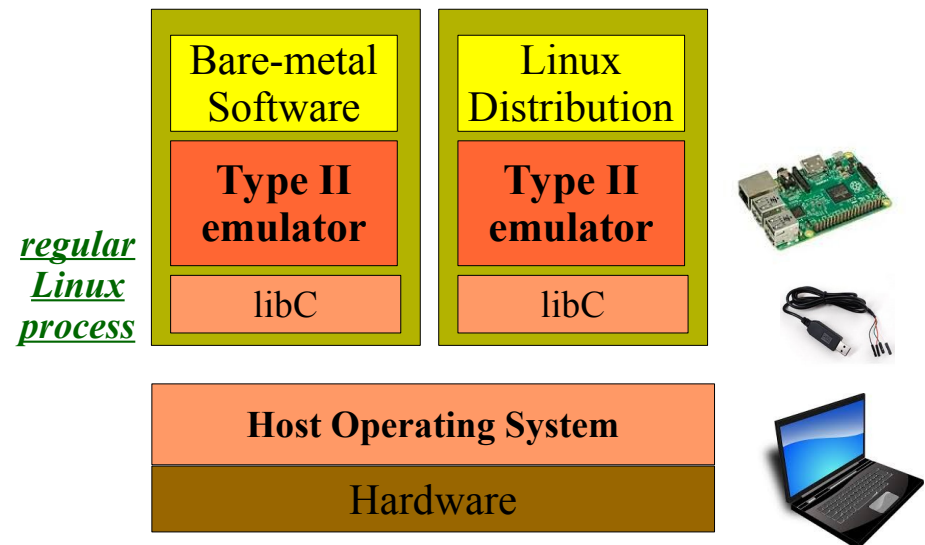
## **Type-II Virtual Machines:**

*Qemu/KVM, VirtualBox, ...*

*Installed as a regular application.*

*Host a single guest operating system, para-virtualized or not.*

*May be used to host a bare-metal application designed for embedded systems.*



(1) Hypervisor, just another name for a virtual machine, there are many kinds of virtual machines...

# QEMU – Booting an VersatilePB Board

- VersatilePB board

- Small ARM-based board
- ARM processor: ARM926EJ-S
- No disk, no display, but a serial line

Requires an arm-none-eabi toolchain

Requires to know the UART controller and its mmio registers

- QEMU

- Kernel: a binary image to load in memory
- Serial: connected to the stdin/stdout of the shell used to launch QEMU

Bare-Metal Software



**Versatile-PB Board**

QEMU

```
$ qemu-system-arm -M versatilepb -serial mon:stdio -kernel your-code
```

UART0:  
Base Address: 0x101F1000

Data register: 0x00  
Flag register: 0x18

*But, how do we know?*



# QEMU – Emulated Hardware

- QEMU Emulated Hardware

- Information via the QEMU monitor
- Example with a PC-104 board

```
$ qemu-system-i386 -serial mon:stdio
```

```
Crtl-A c  
(qemu) info qtree
```

```
...
```

```
dev: i440FX-pcihost, id ""
```

```
  irq 0
```

```
  bus: pci.0
```

```
    dev: PIIX3, id ""
```

```
      class ISA bridge, addr 00:01.0,  
        pci id 8086:7000 (sub 1af4:1100)
```

```
    bus: isa.0
```

```
    type ISA
```

```
      dev: isa-serial, id ""
```

```
        index = 0 (0)
```

```
        iobase = 1016 (0x3f8)
```

```
        irq = 4 (0x4)
```

```
        chardev = "serial0"
```

```
        wakeup = 0 (0)
```

```
        isa irq 4
```

Software

**Linux Kernel**



**PC-104 Board**

**QEMU**

**UART: COM1**

**Base Address: 0x3F8**

**Status: 0x3F8 + 0x05**

**Bit 0x20 → can write a character**

**Bit 0x01 → can read a character**

**Out register: 0x3F8**

**In register: 0x3F8**

**IRQ: 4**

# What's next?

- Hands-on Learning
  - The worklog in Step0...
- Software Installation for Ubuntu/Debian
  - Step0: support the ARM processor toolchain

```
$ sudo apt-get install qemu-system-arm qemu-system-x86
```

```
$ sudo apt-get install gcc-arm-none-eabi gdb-arm-none-eabi binutils-arm-none-eabi
```

Note: if you cannot install the *gdb-arm-none-eabi*, go for the gdb supporting multiple architectures

```
$ sudo apt-get install gdb-multiarch
```

- Step1: support for 32-bit code

```
$ sudo apt-get update -y
```

```
$ sudo apt-get install -y multiarch-support
```

```
$ sudo apt-get install libncurses5:i386
```